


Double-click (or enter) to edit

EDA

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline
plt.style.use('ggplot')
pd.set_option('max_columns',200)

df = pd.read_csv("music_genre.csv")
```

```
# top 5 row in the dataset
df.head()
```



	instance_id	artist_name	track_name	popularity	acousticness	danceability	duration_ms
0	32894.0	Röyksopp	Röyksopp's Night Out	27.0	0.00468	0.652	
1	46652.0	Thievery Corporation	The Shining Path	31.0	0.01270	0.622	2186
2	30097.0	Dillon Francis	Hurricane	28.0	0.00306	0.620	2156
3	62177.0	Dubloadz	Nitro	34.0	0.02540	0.774	1668
4	24907.0	What So Not	Divide & Conquer	32.0	0.00465	0.638	2226

```
#how many number of columns and row
df.shape
```

```
(50005, 18)
```

```
#All columns name
df.columns
```

```
Index(['instance_id', 'artist_name', 'track_name', 'popularity',
       'acousticness', 'danceability', 'duration_ms', 'energy',
       'instrumentalness', 'key', 'liveness', 'loudness', 'mode',
       'speechiness', 'tempo', 'obtained_date', 'valence', 'music_genre'],
      dtype='object')
```

```
#All columns name
df.columns
```

```
Index(['instance_id', 'artist_name', 'track_name', 'popularity',
       'acousticness', 'danceability', 'duration_ms', 'energy',
       'instrumentalness', 'key', 'liveness', 'loudness', 'mode',
       'speechiness', 'tempo', 'obtained_date', 'valence', 'music_genre'],
      dtype='object')
```


```
df.describe()
```

	instance_id	popularity	acousticness	danceability	duration_ms	energy
count	50000.000000	50000.000000	50000.000000	50000.000000	5.000000e+04	50000.000000
mean	55888.396360	44.220420	0.306383	0.558241	2.212526e+05	0.599755
std	20725.256253	15.542008	0.341340	0.178632	1.286720e+05	0.264555

```
sns.set(rc={'figure.figsize':(30,6)})

# Create box plot
sns.boxplot(data=df)

# Show plot
plt.show()
```



```
#Check for missing value
df.isna().sum()
```

instance_id	5
artist_name	5
track_name	5
popularity	5
acousticness	5
danceability	5
duration_ms	5
energy	5
instrumentalness	5
key	5
liveness	5
loudness	5
mode	5
speechiness	5
tempo	5
obtained_date	5
valence	5
music_genre	5
dtype:	int64

```
df= df.dropna()

# from sklearn.preprocessing import LabelEncoder

# le = LabelEncoder()
# df['GENDER'] = le.fit_transform(df['GENDER'])
# df['LUNG_CANCER'] = le.fit_transform(df['LUNG_CANCER'])
```

```
df.dtypes
```

instance_id	float64
artist_name	object
track_name	object
popularity	float64
acousticness	float64
danceability	float64
duration_ms	float64
energy	float64

```

instrumentalness    float64
key                 object
liveness            float64
loudness            float64
mode                object
speechiness         float64
tempo               object
obtained_date       object
valence             float64
music_genre         object
dtype: object

```

```
df.head()
```

	instance_id	artist_name	track_name	popularity	acousticness	danceability	duration_ms
0	32894.0	Röyksopp	Röyksopp's Night Out	27.0	0.00468	0.652	
1	46652.0	Thievery Corporation	The Shining Path	31.0	0.01270	0.622	218000
2	30097.0	Dillon Francis	Hurricane	28.0	0.00306	0.620	215000
3	62177.0	Dubloadz	Nitro	34.0	0.02540	0.774	166000
4	24907.0	What So Not	Divide & Conquer	32.0	0.00465	0.638	222000

```

#convert categorical columns to numbers
# df = port
from sklearn.preprocessing import LabelEncoder

# Create a label encoder object
le = LabelEncoder()

# Iterate over all the values of each column and extract their dtypes
for col in df:
    if df[col].dtype=='object':
        # Use the label encoder object to fit_transform
        df[col]=le.fit_transform(df[col])

```

```
df.head()
```

	instance_id	artist_name	track_name	popularity	acousticness	danceability	duration_ms
0	32894.0	5029	28371	27.0	0.00468	0.652	
1	46652.0	6117	34817	31.0	0.01270	0.622	218000
2	30097.0	1591	15024	28.0	0.00306	0.620	215000
3	62177.0	1707	23372	34.0	0.02540	0.774	166000
4	24907.0	6519	8649	32.0	0.00465	0.638	222000

```
df.dtypes
```

```

instance_id        float64
artist_name        int64
track_name         int64
popularity          float64
acousticness        float64
danceability        float64
duration_ms        float64
energy              float64
instrumentalness    float64
key                 int64
liveness            float64
loudness            float64
mode                int64
speechiness         float64
tempo               int64
obtained_date       int64
valence             float64

```


```
music_genre      int64
dtype: object
```

```
df = df.dropna()
```

```
df_corr = df[['instance_id', 'artist_name', 'track_name', 'popularity',
              'acousticness', 'danceability', 'duration_ms', 'energy',
              'instrumentalness', 'key', 'liveness', 'loudness', 'mode',
              'speechiness', 'tempo', 'obtained_date', 'valence', 'music_genre']].corr()
df_corr
```

	instance_id	artist_name	track_name	popularity	acousticness	danceability
instance_id	1.000000	0.007326	-0.008042	0.006317	-0.005268	0.000372
artist_name	0.007326	1.000000	0.034748	-0.004463	-0.043094	0.000372
track_name	-0.008042	0.034748	1.000000	-0.069139	0.027580	-0.041728
popularity	0.006317	-0.004463	-0.069139	1.000000	-0.290453	0.356420
acousticness	-0.005268	-0.043094	0.027580	-0.290453	1.000000	-0.347681
danceability	0.000372	0.010550	-0.041728	0.356420	-0.347681	1.000000
duration_ms	-0.005848	-0.004754	0.013115	-0.083809	0.061862	-0.109220
energy	0.009952	0.036248	-0.016984	0.216345	-0.791250	0.061862
instrumentalness	-0.004015	-0.019471	0.021728	-0.364960	0.387970	-0.037436
key	0.001189	-0.004723	0.000232	-0.005212	0.005796	-0.037436
liveness	0.004737	0.008216	-0.009292	-0.039468	-0.109220	-0.037436
loudness	0.006847	0.037436	-0.017243	0.317941	-0.730401	0.006189
mode	0.006189	-0.004415	0.003666	0.013427	-0.025161	0.006189
speechiness	-0.005908	-0.007084	-0.030364	0.224309	-0.163377	0.001303
tempo	0.001303	-0.023968	0.001585	-0.012906	0.076434	-0.012906
obtained_date	0.002523	0.018221	-0.008815	0.120444	-0.205964	0.002523
valence	0.005385	0.021287	-0.007354	0.124913	-0.270238	0.005385
music_genre	-0.000861	-0.025123	-0.056281	0.502133	-0.097969	0.000861

```
plt.figure(figsize = (16,10), dpi=200)
ax = plt.axes()
sns.heatmap(df_corr(), annot = True, cmap='RdBu', ax=ax)
ax.set_title('Correlation Matrix - Before Encoding & Handling Missing Data', weight='bold')
plt.show()
```



```
columns = ['key', 'mode', 'tempo']
```

```
for column in columns:
```

```
    print(df[column].unique())
```

```
[ 1  5 11  4  9  2 10  8  0  3  7  6]
[1 0]
[ 274  4193  8138 ... 3631 16869 29157]
```

```
def preprocess_inputs(df):
```

```
    df = df.copy()
```

```
    df = df.drop(['instance_id', 'artist_name', 'track_name', 'obtained_date'], axis=1)
```

```
    df['mode'] = df['mode'].replace({'Minor' : 0,
                                     'Major' : 1})
```

```
    embarked_dummies = pd.get_dummies(df.key)
    df = pd.concat([df, embarked_dummies], axis=1)
    df = df.drop('key', axis=1)
```

```
    df['tempo'] = df['tempo'].replace('?', np.nan)
    df["tempo"] = df["tempo"].astype("float")
    df['tempo'] = df['tempo'].fillna(df['tempo'].mean())
```

```
    df['music_genre'] = df['music_genre'].replace({'Electronic':0, 'Anime':1, 'Jazz':2, 'Alternative':3, 'Country':4, 'Rap':5,
                                                    'Blues':5, 'Rock':6, 'Classical':7, 'Hip-Hop':8})
```

```
    y = df['music_genre']
    X = df.drop('music_genre', axis=1)
```

```
    X_train, X_test, y_train, y_test = train_test_split(X, y, train_size = 0.7, shuffle=True, random_state=43)
```

```
    scaler = StandardScaler()
    scaler.fit(X_train)
```

```
    X_train = pd.DataFrame(scaler.transform(X_train), index = X_train.index, columns = X_train.columns)
    X_test = pd.DataFrame(scaler.transform(X_test), index = X_test.index, columns = X_test.columns)
```

```
    return X_train, X_test, y_train, y_test
```

DT

```
import pandas as pd
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report
```

```

from sklearn.model_selection import train_test_split
from sklearn.tree import DecisionTreeClassifier
from sklearn.neural_network import MLPClassifier
from sklearn.metrics import accuracy_score

# Split the data into features and labels
X = df.drop('music_genre', axis=1)
y = df['music_genre']

# Split the data into training and test sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2)
# train, test, labels_train, labels_test = sklearn.model_selection.train_test_split(iris.data, iris.target, train_size=0.80)

# Define the Decision Tree model
dt = DecisionTreeClassifier()

# Train the Decision Tree model
dt.fit(X_train, y_train)

# Make predictions on the test set
y_pred_dt = dt.predict(X_test)

# Evaluation using accuracy score
acc_dt = accuracy_score(y_test, y_pred_dt)
print("Decision Tree Accuracy:", acc_dt)

# Evaluation using confusion matrix
confusion_matrix_dt = confusion_matrix(y_test, y_pred_dt)
print("Decision Tree Confusion Matrix:")
print(confusion_matrix_dt)

# Evaluation using classification report
print("Decision Tree classification report:")
print(classification_report(y_test, y_pred_dt))

```

Decision Tree Accuracy: 0.4617

Decision Tree Confusion Matrix:

```

[[303  22  54  23 132  84  80  80  78 183]
 [ 29 737  87  62  22  67   0  28   1   6]
 [ 58  72 435  35  72 107 16 147   2  34]
 [ 17  58  37 786  10  28   0  86   0   7]
[128 16  84   9 426  47  29  78  31 136]
 [ 67  64 105  20 40 483 23 142  22  30]
 [ 89   2   4   1 23 28 333 17 441  47]
 [ 74  35 147  67 61 145 20 393 17  24]
 [ 73   1   4   0 31 20 400 12 333 100]
[180   5  46   6 139 20  62  35 112 388]]

```

Decision Tree classification report:

	precision	recall	f1-score	support
0	0.30	0.29	0.29	1039
1	0.73	0.71	0.72	1039
2	0.43	0.44	0.44	978
3	0.78	0.76	0.77	1029
4	0.45	0.43	0.44	984
5	0.47	0.48	0.48	996
6	0.35	0.34	0.34	985
7	0.39	0.40	0.39	983
8	0.32	0.34	0.33	974
9	0.41	0.39	0.40	993
accuracy			0.46	10000
macro avg	0.46	0.46	0.46	10000
weighted avg	0.46	0.46	0.46	10000

```

# input_data = np.array([1,69,1,2,2,1,1,2,1,2,2,2,2,2,2]).reshape(1,-1)
# prediction = dt.predict(input_data)

```

```

# prediction

```

```

!pip install lime
import lime
import lime.lime_tabular
from __future__ import print_function
np.random.seed(1)

```

```

Looking in indexes: https://pypi.org/simple, https://us-python.pkg.dev/colab-wheels/public/simple/
Requirement already satisfied: lime in /usr/local/lib/python3.8/dist-packages (0.2.0.1)
Requirement already satisfied: numpy in /usr/local/lib/python3.8/dist-packages (from lime) (1.21.6)
Requirement already satisfied: scikit-image>=0.12 in /usr/local/lib/python3.8/dist-packages (from lime) (0.18.3)
Requirement already satisfied: matplotlib in /usr/local/lib/python3.8/dist-packages (from lime) (3.2.2)
Requirement already satisfied: tqdm in /usr/local/lib/python3.8/dist-packages (from lime) (4.64.1)
Requirement already satisfied: scipy in /usr/local/lib/python3.8/dist-packages (from lime) (1.7.3)
Requirement already satisfied: scikit-learn>=0.18 in /usr/local/lib/python3.8/dist-packages (from lime) (1.0.2)
Requirement already satisfied: networkx>=2.0 in /usr/local/lib/python3.8/dist-packages (from scikit-image>=0.12->lime) (3.0)
Requirement already satisfied: imageio>=2.3.0 in /usr/local/lib/python3.8/dist-packages (from scikit-image>=0.12->lime) (2.9.0)
Requirement already satisfied: pillow!=7.1.0,!>=7.1.1,>=4.3.0 in /usr/local/lib/python3.8/dist-packages (from scikit-image>=0.12->lime) (
Requirement already satisfied: tifffile>=2019.7.26 in /usr/local/lib/python3.8/dist-packages (from scikit-image>=0.12->lime) (2022.10.16)
Requirement already satisfied: PyWavelets>=1.1.1 in /usr/local/lib/python3.8/dist-packages (from scikit-image>=0.12->lime) (1.4.1)
Requirement already satisfied: kiwisolver>=1.0.1 in /usr/local/lib/python3.8/dist-packages (from matplotlib->lime) (1.4.4)
Requirement already satisfied: pyparsing!=2.0.4,!>=2.1.2,!>=2.1.6,>=2.0.1 in /usr/local/lib/python3.8/dist-packages (from matplotlib->lime)
Requirement already satisfied: cycler>=0.10 in /usr/local/lib/python3.8/dist-packages (from matplotlib->lime) (0.11.0)
Requirement already satisfied: python-dateutil>=2.1 in /usr/local/lib/python3.8/dist-packages (from matplotlib->lime) (2.8.2)
Requirement already satisfied: threadpoolctl>=2.0.0 in /usr/local/lib/python3.8/dist-packages (from scikit-learn>=0.18->lime) (3.1.0)
Requirement already satisfied: joblib>=0.11 in /usr/local/lib/python3.8/dist-packages (from scikit-learn>=0.18->lime) (1.2.0)
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.8/dist-packages (from python-dateutil>=2.1->matplotlib->lime) (1.15.0)

```

```
df.columns
```

```

X=df[['instance_id', 'artist_name', 'track_name', 'popularity',
      'acousticness', 'danceability', 'duration_ms', 'energy',
      'instrumentalness', 'key', 'liveness', 'loudness', 'mode',
      'speechiness', 'tempo', 'obtained_date', 'valence']]

```

```
explainer = lime.lime_tabular.LimeTabularExplainer(X.values, feature_names=df.columns.values.tolist(), class_names=['NO LUNG CANCER', 'LUNG CA
```

```
j = 22
```

```
exp = explainer.explain_instance(X.values[j], dt.predict_proba, num_features=15, top_labels=1)
```

```
exp
```

```

Intercept 0.014863179653321151
Prediction_local [0.42821651]
Right: 1.0
/usr/local/lib/python3.8/dist-packages/sklearn/base.py:450: UserWarning: X does not have valid feature names, but DecisionTreeClassifier
  warnings.warn(
<lime.explanation.Explanation at 0x7f68b3eff190>


```

```
exp.show_in_notebook(show_table=True)
```

Prediction probabilities

NO LUNG C...	<input type="text" value="0.00"/>
LUNG CANCER	<input type="text" value="0.00"/>
Other	<input type="text" value="0.00"/>
	<input type="text" value="0.00"/>
	<input type="text" value="0.00"/>

NO LUNG CANCER LUNG CANCER



▼ Hybrid Model

```
# Use the decision tree to select the most important features from the training data
important_features = dt.feature_importances_
X_train_important = X_train.iloc[:, important_features > 0]
X_test_important = X_test.iloc[:, important_features > 0]

# Train a neural network on the selected features
nn = MLPClassifier()
nn.fit(X_train_important, y_train)

# Make predictions on the test set using the neural network
y_pred_nn = nn.predict(X_test_important)

# Calculate the accuracy of the hybrid model
accuracy = accuracy_score(y_test, y_pred_nn)
print("Accuracy:", accuracy)

# Evaluation using accuracy score
acc_nn = accuracy_score(y_test, y_pred_nn)
print("Hybrid Model Accuracy:", acc_nn)

# Evaluation using confusion matrix
confusion_matrix_nn = confusion_matrix(y_test, y_pred_nn)
print("Hybrid Model Confusion Matrix:")
print(confusion_matrix_nn)

# Evaluation using classification report
print("Hybrid Model report:")
print(classification_report(y_test, y_pred_nn))
```

```
Accuracy: 0.2148
Hybrid Model Accuracy: 0.2148
Hybrid Model Confusion Matrix:
[[ 6  4  0  2  6  0 389 568  0 64]
 [ 6 161  4 79  7  3 39 735  0  5]
 [ 2  15  0 10  5  0 59 868  0 19]
 [ 3  26  1 429  5  0 11 552  0  2]
 [ 4  12  0  1 12  0 261 624  0 70]
 [ 9  13  0  3  6  0 140 807  1 17]
 [ 8  0  2  0  3  0 582 269  0 121]
 [ 5  20  1 16 10  1 66 851  1 12]
 [11  1  0  1  3  0 611 229  2 116]
 [12  1  0  0  5  0 506 364  0 105]]
Hybrid Model report:
              precision    recall  f1-score   support

     0           0.09       0.01       0.01       1039
     1           0.64       0.15       0.25       1039
     2           0.00       0.00       0.00        978
     3           0.79       0.42       0.55       1029
     4           0.19       0.01       0.02        984
     5           0.00       0.00       0.00        996
     6           0.22       0.59       0.32        985
     7           0.15       0.87       0.25        983
     8           0.50       0.00       0.00        974
     9           0.20       0.11       0.14        993

 accuracy          0.21       10000
 macro avg         0.28       0.22       0.15       10000
 weighted avg      0.28       0.21       0.16       10000
```

```
explainer = lime.lime_tabular.LimeTabularExplainer(X_train_important.values, feature_names=df.columns.values.tolist(), class_names=['NO LUNG', 'LUNG'],
# Select a random sample from the test set to explain
i = 21
exp = explainer.explain_instance(X_train_important.iloc[i, :], nn.predict_proba, num_features=14)
```



```


Intercept 0.03591413916701211
Prediction_local [0.05568912]
Right: 2.9774491742121863e-83
/usr/local/lib/python3.8/dist-packages/sklearn/base.py:450: UserWarning: X does not have valid feature names, but MLPClassifier was fitted
  warnings.warn(

```

```

# Show the explanation
exp.show_in_notebook(show_table=True, show_all=False)

```



```

import xgboost as xgb
from sklearn.metrics import accuracy_score

# Split data into training and test sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Create XGBoost data matrices
dtrain = xgb.DMatrix(data=X_train, label=y_train)
dtest = xgb.DMatrix(data=X_test, label=y_test)

# Set XGBoost parameters
params = {'objective': 'binary:logistic', 'eval_metric': 'error'}

# Train the model
model = xgb.train(params=params, dtrain=dtrain, num_boost_round=100)

# Predict on the test set
y_pred = model.predict(dtest)

# Convert predictions to binary labels
y_pred = [round(pred) for pred in y_pred]

# Calculate accuracy
acc = accuracy_score(y_test, y_pred)

# Print accuracy
print("Accuracy: ", acc)

```

```

-----
XGBoostError                                Traceback (most recent call last)
<ipython-input-118-3cd21813e6b7> in <module>
    13
    14 # Train the model
--> 15 model = xgb.train(params=params, dtrain=dtrain, num_boost_round=100)
    16
    17 # Predict on the test set

```

↕ 3 frames

```

/usr/local/lib/python3.8/dist-packages/xgboost/core.py in _check_call(ret)
    174     """
    175     if ret != 0:
--> 176         raise XGBoostError(py_str(_LIB.XGBGetLastError()))
    177
    178

```

XGBoostError: [13:37:22] /workspace/src/objective/regression_obj.cu:101: label must be in [0,1] for logistic regression

Stack trace:

```

[bt] (0) /usr/local/lib/python3.8/dist-
packages/xgboost/.lib/libxgboost.so(dmlc::LogMessageFatal::~LogMessageFatal()+0x24)
[0x7f68aab35cb4]
[bt] (1) /usr/local/lib/python3.8/dist-
packages/xgboost/.lib/libxgboost.so(xgboost::obj::RegLossObj<xgboost::obj::LogisticClas
const&, xgboost::MetaInfo const&, int,
xgboost::HostDeviceVector<xgboost::detail::GradientPairInternal<float> >*>+0x805)
[0x7f68aad3f9d5]
[bt] (2) /usr/local/lib/python3.8/dist-
packages/xgboost/.lib/libxgboost.so(xgboost::LearnerImpl::UpdateOneIter(int,
xgboost::DMatrix*)+0x345) [0x7f68aabc505]
[bt] (3) /usr/local/lib/python3.8/dist-
packages/xgboost/.lib/libxgboost.so(XGBoosterUpdateOneIter+0x35) [0x7f68aab32aa5]

```